

1. Preliminaries

Before starting on this assignment, please be sure to read the General Instructions that are on Canvas (under Files-->General Resources), and also make sure you are able to log in to the class PostgreSQL server. You'll get help on this in your Lab Section, not the Lectures, so *be sure to attend Lab Sections*.

2. Goal

The goal of the first assignment is to create a PostgreSQL data schema with 7 tables. That is all that is required in this assignment; don't do anything that's not required. The other Lab Assignments are much more difficult. In your Lab Sections, you may be given information about how to load data into a table and issue simple SQL queries, because that's fun, but loading data and issuing queries are **not** required in this assignment. (That will show up in the Lab2 assignment.)

3. Lab1 Description

3.1 Create PostgreSQL Schema Lab1

As we noted in the general instructions, you will create a Lab1 schema to set apart the database tables created in this lab from tables you will create in future labs, as well as from tables (and other objects) in the default (public) schema. Note that the meaning of schema here is specific to PostgreSQL, and distinct from the general meaning of schema. See [here](#) for more details on PostgreSQL schemas. You create the Lab1 schema using the following command:

```
CREATE SCHEMA Lab1;
```

[PostgreSQL makes all identifiers lowercase, unless you put them in quotation marks, e.g., "Lab1". But in CSE 180, you don't have to bother using quotation marks for identifiers. We use capitals in assignments for readability, but it's okay (and equivalent) if you use lab1 as the schema name.]

Now that you have created the schema, you want to make Lab1 be the default schema when you use psql. If you do not set Lab1 as the default schema, then you will have to qualify your table names with the schema name (e.g., by writing Lab1.customer, rather than just customer). To set the default schema, you modify your search path as follows. (For more details, see [here](#).)

```
ALTER ROLE username SET SEARCH_PATH to Lab1;
```

You will need to log out and log back in to the server for this default schema change to take effect. (Students **often forget** to do this, and then are surprised that their tables aren't in the expected schema.) To see your current SEARCH_PATH, enter:

```
SHOW SEARCH_PATH;
```

3.2 Tables

You'll be creating tables for a very simplified version of a Conference Management database schema, with tables for Persons, Conferences, Submissions, Authors, Reviewers, Reviews and Attendees. The data types and Referential Integrity for the attributes in these 7 Conference Management tables are described in the next section. No, this schema doesn't provide everything that a real-world Conference Management database would have, but it's a decent start.

Important: To receive full credit, you must use the attribute names as given, and the attributes must be in the order given. Also, the data types and referential integrity must match the specifications given in the next section. Follow directions; do not do more than you're asked to do in this assignment.

Persons(personID, lastName, firstName, email, affiliation, isStudent)

Conferences(conferenceName, year, conferenceDate, regularAttendeeCost, studentAttendeeCost, submissionDueDate, reviewDueDate, importance)

Submissions(conferenceName, year, submissionID, numPages, submitDate, wasAccepted)

Authors(authorID, conferenceName, year, submissionID, authorPosition)

Reviewers(reviewerID, conferenceName, year, reliability)

Reviews(reviewerID, conferenceName, year, submissionID, reviewDate, rating)

Attendees(attendeeID, conferenceName, year)

The underlined attribute (or attributes) identifies the Primary Key of each table. A table can only have one Primary Key, but that Primary Key may involve multiple attributes.

- A person in Persons specifies the person's ID, last name, first name, email, affiliation and whether or not the person is a student. Authors, Reviewers and Attendees are all Persons.
- A conference in Conferences specifies the conference name, year, date, the cost for regular attendees, the cost for student attendees, the due date for papers submitted to the conference, , the due date for reviews of submitted papers, and the importance of the conference (which we'll explain later). For example, the SIGMOD 2021 conference might have SIGMOD as its conferenceName and 2021 as its year; SIGMOD 2022 would have the same conferenceName but a different year.
- A submission in Submissions specifies the conference for which the paper was submitted (by providing conference name and year), the submission's ID, the number of pages in the submission, the date that the paper was submitted, and whether the submission was accepted for the conference.
 - Any (conferenceName, year) that's in a Submissions row **must appear as a** (conferenceName, year) in the Conferences table.

- An author in Authors is a person who submitted a paper to a conference. An author specifies the person who's the author, the submission by that person (by providing conference name and year and submission ID), and the position of the author in the list of authors. For example, if a paper has 4 authors, then you'd expect to have one author in position 1, another in position 2, etc. A person could potentially be an author of multiple submissions for the same or different conferences.
 - Any authorID that's in an Authors row **must appear as** a personID in the Persons table.
 - Any (conferenceName, year) that's in an Authors row **must appear as** a (conferenceName, year) in the Conferences table.
- A reviewer in Reviewers is a person who writes reviews for papers submitted to a conference.. Each reviewer specifies the person who's the reviewer, the conference for which they're a reviewer, and the reliability of the reviewer (which we'll explain later).
 - Any reviewerID that's in a Reviewers row **must appear as** a personID in the Persons table.
 - Any (conferenceName, year) that's in an Reviewers row **must appear as** a (conferenceName, year) in the Conferences table.
- A review in Reviews indicates that a particular reviewer reviewed a particular conference submission on a particular date and gave the submission a particular rating.
 - Any (reviewerID, conferenceName, Year) that's in a Reviews row **must appear as** a (reviewerID, conferenceName, Year) in the Reviewers table.
 - Any (conferenceName, year, submissionID) that's in an Reviews row **must appear as** a (conferenceName, year, submissionID) in the Submissions table.
- An attendee in Attendees is a person who attended a particular conference.
 - Any attendeeID that's in an Attendees row **must appear as** a personID in the Persons table.
 - Any (conferenceName, year) that's in an Attendees row **must appear as** a (conferenceName, year) in the Conferences table.

In this assignment, you'll just have to create tables with the correct table names, attributes, data types, Primary Keys and Referential Integrity. **"Must appear as"** means that there's a Referential Integrity requirement. **Be sure not to forget Primary Keys and Referential Integrity when you do Lab1!**

3.2.1 Data types

Sometimes an attribute (such as year) appears in more than one table. Attributes that have the same attribute name might not have the same data type in all tables, but in our schema, they do.

- For personID, year, authorID, submissionID, numPages, authorPosition, reviewerID, rating and attendeeID, use *integer*.
- For importance (which classifies Conferences) and reliability (which classifies Reviewers), use *character with fixed length 1*. (We'll explain the values of these attributes in a later Lab Assignment.)
- For lastName and firstName, use *character of variable length*, with maximum length 30.
- For conferenceName, email and affiliation, use *character of variable length*, with maximum length 60.
- regularAttendeeCost and studentAttendeeCost should be *numeric*, with at most 4 decimal digits to the left of the decimal point and 2 decimal digits after it.
- conferenceDate, submissionDueDate, reviewDueDate, submitDate and reviewDate should be of type *date*.
- The isStudent attribute, which indicates whether a person is a student, and the wasAccepted attribute, which indicates whether a submission was accepted, should be *boolean*.

You must write a CREATE TABLE statement for each of the 7 tables in Section 3.2. Write the statements in the same order that the tables are listed above. **Use the data types, Primary Keys and Referential Integrity described above.** You will lose credit if you do anything beyond that, even if you think that it's sensible. Save your statements in the file create.sql

Note that PostgreSQL maps all SQL identifiers (e.g., table names and attributes) to lowercase. That's okay in your CSE 180 assignments. You won't lose points for Lab1 because Persons appears in the database as persons, and conferenceName appears as conferencename. It is possible to specify case for identifiers by putting that identifier inside double-quote symbols, e.g., as "Persons". But then every time you refer to that identifier, you'll have to use the double-quotes. "PERSONS" is not the same identifier as "Persons", and neither is the same as Persons, which PostgreSQL maps to persons.

4. Testing

While you're working on your solution, it is a good idea to drop all objects from the schema every time you run the create.sql script, so you can start fresh. Dropping each object in a schema may be tedious, and sometimes there may be a particular order in which objects must be dropped. The following command, which you should put at the top of your create.sql, will drop your Lab1 schema (and all the objects within it), and then create the (empty) schema again:

```
DROP SCHEMA Lab1 CASCADE;  
CREATE SCHEMA Lab1;
```

Before you submit your Lab1 solution, login to your database via psql and execute your create.sql script. As you'll learn in Lab Sections, the command to execute a script is: `\i <filename>`. Verify that every table has been created by using the command: `\d`. Note that when you execute `\d`, the tables that are displayed may appear in any order, not necessarily in the order in which you created them.

To verify that the attributes of each table are in the correct order, and that each attribute is assigned its correct data type use the following command: `\d <table>`.

We'll supply some load data that you can use to test your solution. If your solution fails on that test data, then it's likely that your solution has an error. But although testing can demonstrate that a program is buggy, testing cannot prove that a program is correct.

You do not have to develop your solution on unix.ucsc.edu, but please recognize that we'll run your solution on unix.ucsc.edu. If your solution fails on unix.ucsc.edu, you'll receive a poor grade, even if your solution worked in some other environment.

5. Submitting

1. Save your script as create.sql. You may add informative comments to your scripts if you want. Put any other information for the Graders in a separate README file that you may submit.
2. Zip the file(s) to a single file with name Lab1_XXXXXXX.zip where XXXXXXX is your 7-digit student ID. For example, if a student's ID is 1234567, then the file that this student submits for Lab1 should be named Lab1_1234567.zip

If you have a README file (which is not required), you can use the Unix command:

```
zip Lab1_1234567 create.sql README
```

If you don't have a README file, to create the zip file you can use the Unix command:

```
zip Lab1_1234567 create.sql
```

(Of course, you should use **your own student ID**, not 1234567.) Submit a zip file, even if you only have one file.

Submit the zip file on Canvas under Assignment Lab1. Please be sure that you have access to Canvas for CSE 180. Registered students should automatically have access; students who are not registered in CSE 180 will have Auditor access, so they can't submit solutions. You can replace your original solution, if you like, up to the Lab1 deadline. (Canvas will give the new file a slightly different name, but that's okay.) No students will be admitted to CSE 180 after the Lab1 due date.

If you are working on the UNIX timeshare and your zip file is located there, you will need to copy your file to your computer so that you can upload it to Canvas through your browser. To do that, you will need an FTP (File Transfer Protocol) client to securely transfer files from the UNIX timeshare. A widely used secure FTP client is Filezilla. Installation instructions are found in the site of [FileZilla](#) (make sure you install the distribution suitable for your operating system). After opening the Filezilla client, you will need to set the host field to unix.ucsc.edu, the username to your CruzId and the password to your Blue password, while the port number should be set to 22 (the default port for remote login). By clicking the Quickconnect button, if your credentials are correct, you will connect and be able to see the contents of your remote Unix folder at the right pane (under the title "Remote site"), while the left pane (under the title "Local site") will display the contents of your local file system. With the mouse, you can drag the file from the Unix folder and drop it to the desired location at your computer. This will transfer the file to your local machine, without erasing it from its original remote location. Filezilla is only one of several options for an FTP client. If you are finding it difficult to install the necessary tools and successfully do file transfers, you should promptly ask for help in the Lab Sections; **do not postpone this until the deadline date**. Computers in UCSC Unix Labs also have pre-installed SSH and FTP clients (PuTTY and PSFTP).

Other approaches to copy files includes using SCP (Secure Copy) and using Cut-and-Paste, where you copy the contents of the file from the unix system, and then paste contents into a file on your computer. Cut-and-Paste works well with for small files, but it's a hack that does not work well for large files.

The CSE 180 Teaching Assistants, Dev Purandare and Omkar Ghanekar, will discuss approaches to access unix remotely (SSH for Mac/Linux and PuTTY for Windows) and to move files to your computer (SCP for Mac/Linux and Filezilla for Windows/Mac/Linux) with you during Lab Sections. Attend your Lab Section to ensure that you know how to handle this correctly!

Lab1 is due by 11:59pm on Sunday, October 10. **Late submissions will not be accepted (Canvas won't take them, nor will we), and there will be no make-up Lab assignments.** Always check to make sure that your submission is on Canvas, and that you've submitted the correct file. You will receive **no credit** if you accidentally submit the wrong file, even if you attempt to "prove" that you completed the correct file on time.